

Stefan Breuers <breuers@vision.rwth-aachen.de>
Wolfgang Mehner <mehner@vision.rwth-aachen.de>

Exercise 5: Matching, Homographies

due **before** 2017-01-09

Important information regarding the exercises:

- The exercise is not mandatory.
- There will be no corrections.
- Nevertheless, we encourage you to work on the exercises and present your solutions in the exercise class. For this regard the submission rules.
- In the archive for this exercise you will find the functions `apply.m` that should be used for displaying your results. You should also use it to test your implementation and see if the results make sense. Answers are to be submitted within `answers.m` Do **not** modify the `apply` files in any way.
- Please do **not** include the data files in your submission!
- If applicable submit your code solution as a zip/tar.gz file named `mn1_mn2_mn3.{zip/tar.gz}` with your **matriculation numbers** (mn).
- Submit your solutions via the L²P system.

Question 1: Region Descriptors ($\Sigma = 0$)

In order to find correspondences between interest points, we need to design region descriptors. In this question, we will implement several **simple descriptors** based on the **histogram** representations from exercise sheet 4.

You can test your implementations either using the points you detected in the previous exercise or using the points from the file `ip_graff.mat`. This file contains points detected with the **Harris detector** using the parameter settings `sigma = 2.0, thresh = 100000`. The vectors `px1` and `py1` are point coordinates from the image `graff5/img1.png` and `px2` and `py2` are point coordinates from the image `graff5/img2.png`.

The functions `histr_g` and `histr_dxdy` are modified versions of `myhist3` and `myhist4` from a previous exercise sheet, which take the input image as a first parameter.

- (a) As an example, we provide the function `descriptors_rg` which takes an input image and a list of interest points and computes **an r, g color histogram over** the $m \times m$ sub-windows **around each interest point** (using the function `histr_g`).

Write a similar function `descriptors_dxdy` which computes **a dx, dy histogram** around each interest point (using the function `histr_dxdy`).

Hint: For our application here, suitable descriptor parameters are a window size of $m=41$, a histogram resolution of `bins=16` and `sigma=2.0`.

```
1 function D = descriptors_dxdy(img, px, py, size, sigma, bins)
```

Visualize the histograms for some interest points. Do they look useful?

- (b) Write a another function `histr_maglap` and `descriptors_maglap` which compute a *mag, lap* (**gradient magnitude, laplacian**) histogram around each interest point. For this descriptor, you can also use the parameters from question a.

```

1 function h = histmaglap(img, sigma, bins)
2 function D = descriptors_maglap(img, px, py, size, sigma, bins)

```

Visualize the histograms for some interest points. Do they look useful?

- (c) Given two histograms h_1 and h_2 with D bins, the χ^2 distance is defined as following:

$$\text{dist_chi2}(h_1, h_2) = \sum_{i=1}^D \frac{(h_1(i) - h_2(i))^2}{h_1(i) + h_2(i)}$$

In the nominator it matches the squared Euclidean distance, and in the denominator each cell is normalized by the total sum. This means that the cells are not weighted equally anymore, but bins with less data have a higher impact on the distance. This has a statistical background. The χ^2 distance tests if the underlying distributions of the histograms are different. In general the χ^2 distance is more robust to outliers, given the histograms contain enough observations.

Write a function `findnn` which takes two sets of region descriptors `D1` and `D2` and tries to find for each descriptor in `D1` the nearest neighbor in `D2` using the Euclidean distance (Matlab command `norm`). The function should return the indices `Idx` and distances `Dist` of the nearest neighbors. Create a second function `findnn_chi2` which does the same using the χ^2 distance.

```

1 function [Idx, Dist] = findnn(D1, D2)
2 function [Idx, Dist] = findnn_chi2(D1, D2)

```

Find the best match using `descriptors_rg` and display the matching histograms. Do they look similar enough?

Question 2: Matching ($\Sigma = 0$)

Now we have all the components for a small matching application. In the exercise archive you will find the folders `graff5` and `NewYork` that contain test scenes with controlled image-plane rotations and viewpoint changes, for which we will try to find point correspondences.

- (a) First we want to try out the color descriptors. Load the two example images `graff5/img1.png` and `graff5/img2.png` and perform the following steps.
 1. Compute Harris interest points for both images. (Use a high threshold to limit the number of interest point to less than 1000)
 2. Compute r/g color histogram descriptors for all interest points.
 3. Find the best matches using the functions `findnn` and `findnn_chi2`.
 4. Use the function `displaymatches` to visualize the N best matches. What do you observe?
- (b) Now, try the same with Hessian interest points and the dx/dy histogram descriptors. What do you observe? Which combination gives the better results? Can you think of an explanation?
- (c) Next, we will try to find matches under image plane rotations. The folder `NewYork` contains a series of test images for which the true homographies are known. In the exercise archive, you can find a small Matlab program `hom_gui_H` which lets you visualize the corresponding point locations. Load the two images `NewYork/im1.png` and `NewYork/im5.png` and start the program as follows. This will open a window showing the two images side by side. Click on one image and describe what happens.

```

1 img1 = double(imread('NewYork/im1.png'));
2 img2 = double(imread('NewYork/im5.png'));
3 H = load('NewYork/H1to5');
4 hom_gui_H(uint8(img1), uint8(img2), H);

```

- (d) Try to find matches between the two images NewYork/im1.png and NewYork/im5.png using Hessian interest points and the dx/dy histogram descriptors. What do you observe?

Now try the *mag/lap* histogram descriptors on the same image pair. What performance do you get? Which descriptor performs better? Why do you think that is?

Question 3: Homography Estimation ($\Sigma = 0$)

Since the test images show only planar scenes, we can try to estimate the homography \mathbf{H} which transforms one of the images into the other. In the following, we briefly repeat the procedure introduced in the lecture. In the literature, this procedure is known as the *Direct Linear Transformation* (DLT) algorithm.

For a point \mathbf{x}_r in the reference image and a corresponding point \mathbf{x}_t in the transformed image, the transformation can be written as follows:

$$\mathbf{H}\mathbf{x}_r = \mathbf{x}'_t \quad (1)$$

$$\begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & 1 \end{bmatrix} \begin{bmatrix} x_r \\ y_r \\ 1 \end{bmatrix} = \begin{bmatrix} x'_t \\ y'_t \\ z'_t \end{bmatrix} \quad \text{with} \quad \frac{1}{z'_t} \begin{bmatrix} x'_t \\ y'_t \\ z'_t \end{bmatrix} = \begin{bmatrix} x_t \\ y_t \\ 1 \end{bmatrix}, \quad (2)$$

and we obtain a system of linear equations with 8 unknowns h_{11} , h_{12} , h_{13} , h_{21} , h_{22} , h_{23} , h_{31} , and h_{32} :

$$\frac{h_{11}x_r + h_{12}y_r + h_{13}}{h_{31}x_r + h_{32}y_r + 1} = x_t \quad (3)$$

$$\frac{h_{21}x_r + h_{22}y_r + h_{23}}{h_{31}x_r + h_{32}y_r + 1} = y_t \quad (4)$$

which can be rewritten as follows

$$h_{11}x_r + h_{12}y_r + h_{13} - x_th_{31}x_r - x_th_{32}y_r - x_t = 0 \quad (5)$$

$$h_{21}x_r + h_{22}y_r + h_{23} - y_th_{31}x_r - y_th_{32}y_r - y_t = 0 \quad (6)$$

In order to estimate these 8 parameters, we need at least 4 corresponding point pairs.

However, we can enhance the accuracy of the estimated homography by using more than 4 point pairs $(\mathbf{x}_1, \mathbf{x}'_1), \dots, (\mathbf{x}_n, \mathbf{x}'_n)$. This leads to an overdetermined system of equations

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -x'_1y_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1x_1 & -y_1y'_1 & -y'_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2x'_2 & -x'_2y_2 & -x'_2 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -y'_2x_2 & -y_2y'_2 & -y'_2 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_n & y_n & 1 & 0 & 0 & 0 & -x_nx'_n & -x'_ny_n & -x'_n \\ 0 & 0 & 0 & x_n & y_n & 1 & -y'_nx_n & -y_ny'_n & -y'_n \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \quad (7)$$

$$\mathbf{A}\mathbf{h} = \mathbf{0} \quad (8)$$

which we want to solve by minimizing the least-squares error. If \mathbf{A} is square, we can directly obtain an exact solution. However, if the system is overdetermined (i.e. $n > 4$), the matrix \mathbf{A} is not square. This problem can be solved by building the so-called *pseudo-inverse* $\mathbf{A}^T \mathbf{A}$, which is square and can therefore be decomposed by eigenvalue decomposition. The solution is the (unit) eigenvector of $\mathbf{A}^T \mathbf{A}$ with least eigenvalue (Matlab command **eig**). Equivalently (and computationally more efficiently), the solution can be obtained by SVD as the unit singular vector corresponding to the smallest singular value of \mathbf{A} (Matlab command **svd**).

$$\mathbf{A} \stackrel{\text{svd}}{=} \mathbf{U} \mathbf{D} \mathbf{V}^T = \mathbf{U} \begin{bmatrix} d_{11} & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & d_{99} \end{bmatrix} \begin{bmatrix} v_{11} & \cdots & v_{19} \\ \vdots & \ddots & \vdots \\ v_{91} & \cdots & v_{99} \end{bmatrix}^T \quad (9)$$

In this formulation, the eigenvector of $\mathbf{A}^T \mathbf{A}$ corresponding to the smallest eigenvalue minimizes the least-squares error to the solution for \mathbf{h} . We can therefore obtain the homography \mathbf{h} from the last column of \mathbf{V} (since we require $h_{33} = 1$, we normalize with v_{99}):

$$\mathbf{h} = \frac{[v_{19}, \dots, v_{99}]}{v_{99}} \quad (10)$$

- (a) Write a function `estimate_homography` which approximates the homography between two images from a set of point correspondences according to the procedure described above.
1. Build up the matrix \mathbf{A} according to equation (7).
 2. Apply SVD to decompose the matrix using the Matlab command `[U, S, V] = svd(A)`. (*Caution:* Maybe you noted that the matrix \mathbf{V} is transposed (see Equation (9)). Matlab also returns a matrix that is transposed in the same way.)
 3. Compute \mathbf{h} according to Equation (10).
 4. Transform \mathbf{h} into a 3×3 matrix \mathbf{H} using the **reshape** command and then transpose it so it is row-major.

```
1 function H = estimate_homography(px1,py1,px2,py2)
```

- (b) Compute the best 10 matches between the two images `NewYork/im1.png` and `NewYork/im5.png` using Hessian interest points and the *mag/lap* histogram descriptors (Matlab command **sort**). Estimate the homography \mathbf{H} from the matches and compare it to the ground truth matrix from the file `NewYork/H1to5`. How accurate is your estimate? Does it improve when you take the first 20 or 50 matches instead? Why/why not?
- (c) Try the accuracy of your estimated homography using the demo program from Question 2. What do you observe?